

Aggregate CRD Design Options

OpenAPI Operator Generator

This document outlines design options for adding an aggregate CRD feature to the OpenAPI operator generator. An aggregate CRD takes one or more generated CRDs as input and combines them for unified management or observation.

Option 1: Composition CRD (Reference-Based)

A CRD that references existing Custom Resources by name/namespace and aggregates their statuses.

```
apiVersion: petstore.example.com/v1alpha1
kind: PetStoreAggregate
metadata:
  name: my-store-health
spec:
  resources:
    - kind: Pet
      name: my-pet
      namespace: default
    - kind: Order
      name: my-order
      namespace: default
  aggregationStrategy: AllHealthy # or AnyHealthy, Quorum
status:
  state: Synced # Aggregated state
  resourceStatuses:
    - kind: Pet
      name: my-pet
      state: Synced
      externalID: "123"
    - kind: Order
      name: my-order
      state: Failed
      message: "API returned 404"
  conditions: [...]
```

Pros:

- Clean separation - resources managed independently
- Leverages existing AllHealthy pattern in pkg/endpoint/resolver.go
- Easy to understand ownership model

Cons:

- Requires watching multiple resource types
- Resources must exist before aggregate

Option 2: Inline Composition CRD (Embedded Specs)

An aggregate CRD that embeds the specs of child resources directly, creating multiple API resources from a single CR.

```
apiVersion: petstore.example.com/v1alpha1
kind: PetBundle
metadata:
  name: full-pet-setup
spec:
  pet:
    name: "Fido"
    status: available
  order:
    petId: "{{.pet.externalID}}" # Template reference
    quantity: 1
status:
  pet:
    state: Synced
    externalID: "123"
  order:
    state: Synced
    externalID: "456"
```

Pros:

- Single CR creates multiple API resources
- Can express dependencies between resources (order references pet ID)
- Familiar pattern (similar to Helm charts, Argo Application)

Cons:

- More complex reconciliation logic (ordering, dependency resolution)
- Harder to manage partial failures
- Spec becomes large

Option 3: Declarative Bundle (Template-Based)

A meta-CRD that declares a bundle template, instantiated by a separate CR. Provides reusable templates with parameters.

```
# Template definition
apiVersion: petstore.example.com/v1alpha1
kind: PetBundleTemplate
metadata:
  name: pet-with-order
spec:
  parameters:
    - name: petName
      type: string
    - name: quantity
      type: integer
      default: 1
  resources:
    - apiVersion: petstore.example.com/v1alpha1
      kind: Pet
      spec:
        name: "{{ .petName }}"
    - apiVersion: petstore.example.com/v1alpha1
      kind: Order
      spec:
        petId: "{{ .resources.Pet.status.externalID }}"
        quantity: "{{ .quantity }}"
---
# Instance
apiVersion: petstore.example.com/v1alpha1
kind: PetBundle
metadata:
  name: my-bundle
spec:
  templateRef: pet-with-order
  parameters:
    petName: "Buddy"
    quantity: 2
```

Pros:

- Reusable templates
- Clear separation of template definition and instantiation
- Parameters enable customization

Cons:

- Most complex to implement
- Two new CRD types needed
- Template language adds complexity

Option 4: Status Aggregator (Read-Only)

A lightweight CRD that only observes and aggregates status from existing resources. Similar to the existing ReadOnly mode.

```

apiVersion: petstore.example.com/v1alpha1
kind: ResourceHealthCheck
metadata:
  name: store-health
spec:
  selector:
    matchLabels:
      app: petstore
    # OR explicit list
  resources:
    - group: petstore.example.com
      kind: Pet
    - group: petstore.example.com
      kind: Order
  aggregation:
    type: Summary # or Detailed
status:
  summary:
    total: 10
    synced: 8
    failed: 2
  conditions:
    - type: AllHealthy
      status: "False"
      reason: "2 resources failed"

```

Pros:

- Simplest to implement - read-only aggregation
- Builds on existing ReadOnly mode pattern
- Good for dashboards/monitoring

Cons:

- Cannot create/manage resources
- Limited to status observation

Option 5: Workflow/Pipeline CRD

An ordered sequence of operations with dependencies. Steps execute in order based on dependency graph.

```
apiVersion: petstore.example.com/v1alpha1
kind: PetWorkflow
metadata:
  name: onboard-pet
spec:
  steps:
    - name: create-pet
      resource:
        kind: Pet
      spec:
        name: "Fido"
        status: pending
    - name: upload-image
      dependsOn: [create-pet]
      action:
        kind: PetUploadImage
      spec:
        petId: "{{ steps.create-pet.externalID }}"
        file: "..."
    - name: mark-available
      dependsOn: [upload-image]
      resource:
        kind: Pet
      spec:
        name: "Fido"
        status: available
  status:
    phase: Running # Pending, Running, Succeeded, Failed
    steps:
      create-pet: Completed
      upload-image: Running
      mark-available: Pending
```

Pros:

- Expresses ordered operations naturally
- Good fit for Action endpoints (one-shot operations)
- Clear execution semantics

Cons:

- Most complex reconciliation logic
- Need DAG evaluation
- State machine complexity

Implementation Considerations

Aspect	Current Support	What's Needed
Multi-resource status	Yes (Responses map)	Extend to cross-CRD

Type references	Yes (UsesSharedType)	Cross-CRD references
Health aggregation	Yes (AllHealthy strategy)	Apply to CRs not endpoints
Watches	Single type	Multi-type watches
Templates	Go templates	Add aggregate templates

Recommendation

For a first iteration, **Option 1 (Composition CRD)** or **Option 4 (Status Aggregator)** are recommended:

- **Option 1** - If you need the aggregate to manage lifecycle (create/delete child resources)
- **Option 4** - If you only need observability across resources

Both leverage existing patterns (Responses map, AllHealthy strategy, ReadOnly mode) and require minimal new concepts.

Crossplane function-status-transformer Analysis

Overview

The Crossplane **function-status-transformer** is a composition function that provides a declarative, matcher-based approach to status aggregation. This pattern is directly applicable to the aggregate CRD design.

What It Does:

- Watches composed resources for status condition changes
- Matches conditions using configurable matchers (exact match, regex, wildcards)
- Transforms and propagates matched conditions to parent composite resource
- Emits Kubernetes events based on condition transitions

Key Design Pattern: Matcher-to-Action Pipeline

The function implements a hook-based pattern where matchers select resources and conditions, then actions propagate transformed status to the parent.

```
StatusConditionHooks[]  
|-- Matchers (AND logic between matchers)  
|   |-- resources: ["cloudsql-\d+"]      # regex selection  
|   |-- conditions:                      # what to match  
|       |   type: Synced  
|       |   status: "False"  
|       |   message: "error: (?P<Error>.+)"  # capture groups  
|   +- matchType: AnyResourceMatchesAllConditions  
+- Actions  
    |-- setConditions:                  # propagate to parent  
    |   type: DatabaseReady  
    |   message: "Failed: {{ .Error }}"  # template injection  
    +- createEvents: [...]
```

Match Types

- **AnyResourceMatchesAnyCondition** - Any resource matches any condition
- **AnyResourceMatchesAllConditions** - Any resource matches all conditions
- **AllResourcesMatchAnyCondition** - All resources match any condition
- **AllResourcesMatchAllConditions** - All resources match all conditions (default)

Key Features

- **Regex capture groups** - Extract error details: message: "error: (?P<Error>.+)""
- **Template injection** - Use captured data: "Failed: {{ .Error }}"
- **Sequential processing** - Hooks execute in order, later hooks see earlier state
- **Non-blocking errors** - Failures set condition status, don't halt reconciliation

Application to Aggregate CRD

The function-status-transformer pattern maps well to **Option 4 (Status Aggregator)** with potential extensions to **Option 1 (Composition CRD)**.

```
apiVersion: petstore.example.com/v1alpha1
kind: ResourceAggregate
metadata:
  name: pet-health
spec:
  # Resource selection (similar to Crossplane matchers)
  resourceSelectors:
    - group: petstore.example.com
      kind: Pet
      matchLabels:
        app: petstore
    - group: petstore.example.com
      kind: Order
      namePattern: "order-*" # regex support

  # Status condition hooks (from function-status-transformer)
  statusConditionHooks:
    - matchers:
        - resources:
            kind: Pet
            conditions:
              - type: Synced
                status: "False"
                reason: "APIError"
                message: "(?P<Error>.+)"
              - matchType: AnyResourceMatchesAnyCondition
            setConditions:
              - type: PetsHealthy
                status: "False"
                reason: SyncFailure
                message: "Pet sync failed: {{ .Error }}"

        - matchers:
            - resources:
                kind: "*" # all resources
                conditions:
                  - type: Synced
                    status: "True"
            - matchType: AllResourcesMatchAllConditions
            setConditions:
              - type: Ready
                status: "True"
                reason: AllSynced
                message: "All {{ .MatchedCount }} resources synced"

  status:
    conditions:
      - type: Ready
        status: "False"
        reason: SyncFailure
      - type: PetsHealthy
        status: "False"
        message: "Pet sync failed: API returned 503"
  matchedResources:
    - kind: Pet
      name: my-pet
      state: Failed
    - kind: Order
      name: order-123
```

state: Synced

Pattern Mapping

Crossplane Pattern	Application to Generator
Matcher-based selection	Select resources by kind, name regex, labels
Condition matching	Match on State, ExternalID, Message fields
Regex capture groups	Extract error details for aggregated messages
Match types (Any/All)	AllHealthy, AnyHealthy, Quorum(n) strategies
Template injection	{{ .Error }}, {{ .MatchedCount }} in messages
Non-blocking errors	Set condition to False rather than failing reconciliation
Sequential hooks	Process hooks in order, later hooks see earlier state

Differences from Crossplane

Aspect	Crossplane	OpenAPI Generator
Resource type	Generic K8s resources	Generated CRDs only
Execution model	gRPC function pipeline	In-controller reconciliation
Status source	status.conditions[]	status.state, status.message
Discovery	Composition defines resources	Selector-based discovery

Implementation Sketch

```
// In pkg/mapper/resource.go - new aggregate type
type AggregateCRDDefinition struct {
    Kind           string
    ResourceSelectors []ResourceSelector
    StatusHooks    []StatusConditionHook
}

type ResourceSelector struct {
```

```

Group      string
Kind       string
NamePattern string          // regex
MatchLabels map[string]string
}

type StatusConditionHook struct {
    Matchers     []ConditionMatcher
    MatchType    MatchType // AnyAny, AnyAll, AllAny, AllAll
    SetConditions []ConditionTemplate
}

type ConditionMatcher struct {
    ResourceKind string
    Conditions   []ConditionMatch
}

type ConditionMatch struct {
    Field  string // "State", "Message", "ExternalID"
    Pattern string // regex with capture groups
}

```

Conclusion

The function-status-transformer design is a **strong fit** for the aggregate CRD feature. Recommendations:

1. **Adopt the matcher-to-action pattern** - It's declarative, flexible, and battle-tested
2. **Simplify for your use case** - No gRPC pipelines needed; inline logic in controller
3. **Use existing status fields** - Match on State, Message, ExternalID instead of conditions
4. **Add match types** - AllHealthy, AnyHealthy, Quorum(n) map to Crossplane's match types

Sources:

- github.com/crossplane-contrib/function-status-transformer
- github.com/crossplane/crossplane/blob/main/design/design-doc-composition-functions.md